

```
In [ ]: '''  
        Author: A.Shrikant  
        '''
```

```
In [1]: # Basic Libraries  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
# Importing the datasets and the DNN Libraries  
import keras  
from keras import datasets  
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense  
from keras.models import Sequential
```

```
In [2]: (x_train, y_train), (x_test, y_test) = datasets.mnist.load_data()
```

```
In [3]: # data shape  
print(x_train.shape)  
print(y_train.shape)  
print(x_test.shape)  
print(y_test.shape)
```

```
(60000, 28, 28)  
(60000,)  
(10000, 28, 28)  
(10000,)
```

```
In [4]: x_train_min = x_train.min()  
x_train_max = x_train.max()  
  
print(f'x_train_min: {x_train_min}')  
print(f'x_train_max: {x_train_max}')
```

```
x_train_min: 0  
x_train_max: 255
```

```
In [5]: # Scale the data.  
x_train = x_train/255.0
```

```
x_train_min: 0.0
x_train_max: 1.0
```

```
Out[7]: (28, 28)
```

[0 1 2 3 4 5 6 7 8 9]

```
Out[10]: 'viridis'
```

```
Out[11]: viridis
```



```
In [12]: # plt.imshow(X, cmap=None)


# X is a multidimensional array representing image data.

# When image has shape (M,N,3) or (M,N,4), the values in image are interpreted
# as RGB or RGBA values. In this case the cmap is ignored.
# cmap : str or ~matplotlib.colors.Colormap, default: image.cmap


# https://stackoverflow.com/questions/25625952/matplotlib-what-is-the-function-of-cmap-in-imshow

def show_img(idx):
    plt.figure(figsize=(20, 2))
    plt.imshow(x_train[idx, :], cmap='gray')
    plt.xlabel(f'Actual label: {y_train[idx]}')
    plt.colorbar()
    plt.show()
```


```
In [13]: show_img(0)
```




```
In [14]: show_img(10)
```



```
In [15]: show_img(20)
```



```
In [16]: show_img(40)
```



Building the CNN based model:

```
In [17]: # Model Building

model = Sequential()

# For Conv2D():

# kernel_size is the filter size(k). Use odd number of filters in convolution.
# So that when padding is 'same' the value p comes out as a whole number.

# padding='valid' by default this means no padding
# padding='same' means that the padding is automatically calculated.
# When strides = 1 the output (feature map) has the same spatial dimensions as the input
# image using the formula  $p = (k-1)/2$ .
# When strides > 1 the output (feature map) has the following spatial dimensions:
#  $\text{output\_spatial\_shape}[i] = \text{ceil}(\text{input\_spatial\_shape}[i] / \text{strides}[i])$ 

# strides=(1, 1) by default.

# input_shape=(batchSize, height, width, channels)
# Specifying the batchSize here in input_shape won't have any effect. The
# recommended place to specify the batchSize is the model.fit() method.
model.add(Conv2D(filters=16, kernel_size=(3,3), strides=(1, 1), padding='valid', activation='relu', input_shape=(28, 28, 1)))

# For MaxPooling2D():
```

```

# pool_size=(2, 2) by default
# strides=(2, 2) by default.
# output_shape = math.floor((input_shape - pool_size) / strides) + 1
model.add(MaxPooling2D((2,2)))

model.add(Conv2D(filters=16, kernel_size=(3,3), strides=(1, 1), padding='valid', activation='relu'))

model.add(MaxPooling2D((2,2)))

# Flatten Layer in Keras converts a multi-dimensional array into a 1D array.
model.add(Flatten())

model.add(Dense(units=16, activation='relu'))

model.add(Dense(units=16, activation='relu'))

model.add(Dense(units=10, activation='softmax'))

# Model compilation

# keras.metrics.SparseCategoricalAccuracy(): Calculates how often predictions
# match integer labels.
# When loss='sparse_categorical_crossentropy', 'accuracy' metric works same
# as keras.metrics.SparseCategoricalAccuracy()
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

```

In [18]: *# What does None in the Output Shape column of model.summary() mean?*

*# The first dimension in Output Shape tuple is the batch size. None usually
means that the size of that dimension is not fixed and may vary based on the
input data.*

<https://stackoverflow.com/questions/47240348/what-is-the-meaning-of-the-none-in-model-summary-of-keras>

*# Output shape of the convolution layer = (n-k+2p)/s+1
Number of paramaters associated with the convolution layer = (k*k+1) * #filters*

Output shape of the max-pooling layer = (n-k)/s+1

```

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d (MaxPooling2D)	(None, 13, 13, 16)	0
conv2d_1 (Conv2D)	(None, 11, 11, 16)	2320
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 16)	0
flatten (Flatten)	(None, 400)	0
dense (Dense)	(None, 16)	6416
dense_1 (Dense)	(None, 16)	272
dense_2 (Dense)	(None, 10)	170
=====		
Total params: 9338 (36.48 KB)		
Trainable params: 9338 (36.48 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [19]: history = model.fit(x_train, y_train, epochs=10, validation_split=0.2)
```

```

Epoch 1/10
1500/1500 [=====] - 38s 23ms/step - loss: 0.3685 - accuracy: 0.8830 - val_loss: 0.1363 - va
l_accuracy: 0.9592
Epoch 2/10
1500/1500 [=====] - 26s 18ms/step - loss: 0.1151 - accuracy: 0.9646 - val_loss: 0.0934 - va
l_accuracy: 0.9730
Epoch 3/10
1500/1500 [=====] - 28s 18ms/step - loss: 0.0847 - accuracy: 0.9741 - val_loss: 0.0769 - va
l_accuracy: 0.9778
Epoch 4/10
1500/1500 [=====] - 26s 18ms/step - loss: 0.0671 - accuracy: 0.9794 - val_loss: 0.0718 - va
l_accuracy: 0.9799
Epoch 5/10
1500/1500 [=====] - 26s 17ms/step - loss: 0.0562 - accuracy: 0.9829 - val_loss: 0.0619 - va
l_accuracy: 0.9822
Epoch 6/10
1500/1500 [=====] - 27s 18ms/step - loss: 0.0488 - accuracy: 0.9848 - val_loss: 0.0614 - va
l_accuracy: 0.9828
Epoch 7/10
1500/1500 [=====] - 27s 18ms/step - loss: 0.0430 - accuracy: 0.9864 - val_loss: 0.0611 - va
l_accuracy: 0.9834
Epoch 8/10
1500/1500 [=====] - 27s 18ms/step - loss: 0.0386 - accuracy: 0.9876 - val_loss: 0.0509 - va
l_accuracy: 0.9858
Epoch 9/10
1500/1500 [=====] - 37s 25ms/step - loss: 0.0334 - accuracy: 0.9895 - val_loss: 0.0474 - va
l_accuracy: 0.9861
Epoch 10/10
1500/1500 [=====] - 28s 19ms/step - loss: 0.0306 - accuracy: 0.9902 - val_loss: 0.0521 - va
l_accuracy: 0.9852

```

Evaluating the model:

```

In [20]: # The model.evaluate(x_test, y_test) function uses the trained neural network to
# make predictions on the input data x_test and then compares those predictions
# with the true labels y_test to calculate the loss and any specified metrics.

model.evaluate(x_test, y_test)

```

```

313/313 [=====] - 3s 10ms/step - loss: 0.0432 - accuracy: 0.9862


```


Out[20]: [0.04315326362848282, 0.9861999750137329]

```
In [21]: plt.figure(figsize=(5, 5))


plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.ylabel('accuracy')
plt.xlabel('epochs')
plt.legend(['training', 'testing'])

plt.show()
```




```
In [22]: y_pred = model.predict(x_test)
```

313/313 [=====] - 2s 5ms/step

In [23]: `y_pred.shape`Out[23]: `(10000, 10)`In [24]: `# First 10 data points predictions.
y_pred_class = [np.argmax(element) for element in y_pred]
y_pred_class[:10]`Out[24]: `[7, 2, 1, 0, 4, 1, 4, 9, 5, 9]`In [25]: `y_test[:10]`Out[25]: `array([7, 2, 1, 0, 4, 1, 4, 9, 5, 9], dtype=uint8)`In [26]: `def act_to_pred(val):
 plt.figure(figsize=(20, 2))
 plt.imshow(x_test[val, :])
 plt.xlabel(f' Actual label: {y_test[val]}')
 plt.show()
 print('The predicted image label is:', y_pred_class[val])`In [27]: `act_to_pred(0)`


The predicted image label is: 7

In [28]: `act_to_pred(100)`



The predicted image label is: 6

```
In [29]: act_to_pred(5784)
```



The predicted image label is: 9

Conclusion:

The CNN based model to predict the hand written digit from image has a **train accuracy of 99.02%**, **validation accuracy of 98.52%** and **test accuracy of 98.62%**.